

ETRA Challenge Report

Dávid Kubek

January 31, 2023

In 2019, the [ETRA](#) organization announced a challenge for the analysis of a dataset pertaining to human eye-movement. The objective of the challenge was to utilize various tools to uncover intriguing relationships and information within the data, potentially leading to novel insights. Additional details regarding the challenge can be found in the official [Challenge Track Call for Papers](#) issued by ETRA.

Getting data

Download zip file of the [dataset](#). Description of the dataset and how the data was collected can be found on [ETRA dataset description](#).

The directory `data` contains the following directories/files

- `data/` : Contains subdirectories containing data of individual participants. Each subdirectory is filled with CSV files, each containing 45 seconds of eye tracker data.
- `images/` : Contains the images shown to the participants for each task.
- `DataSummary.csv` : Contains mouse clicks of the participants for tasks like Puzzle or Where's Waldo.

Hypotheses

We will be interested in the following hypotheses:

1. **Is there a point where a participant becomes tired/bored when viewing a scene? (the frequency of saccades decreases) If so, are some scenes less interesting than others?**

There are different types of scenes in the study and some might be more visually stimulating than others. For example, it is not unreasonable to expect that a participant will

stay much more engaged in a scene with a lot of detail to study than when looking at a blank scene. Another possibility is that during visual searches (picture puzzles, where is Waldo) a participant may become frustrated and just give up on the task.

The reasoning behind the way we are going to measure this is as follows. When a participant is engaged in a scene, we assume that their gaze is going to dart across the scene rapidly, as they are studying all the details and the number of saccades during a fixed time frame is going to be very high. On the other hand, when subject becomes bored or frustrated, we might expect that they are going to rest their gaze at some point and not move their eyes very much (maybe just waiting for the trial to be over) and so the frequency of saccades would decrease.

However, the decrease in frequency of saccades might not be exactly indicative of the subjects boredom or frustration. It is not unlikely that the frequency of saccades is high every time a new scene is introduced and naturally decreases as the subject becomes acquainted with the scene.

Whatever the reason might be, if there indeed is such a changing point in the frequency of saccades, it might be interesting to study the time for such change and it's relationship to participant, scene and task.

2. Does the pupil size change depending on scene? What about dependence on task?

The motivation behind this hypothesis is that we would like to study the excitement of a subject viewing some scene or performing some task. One of the indicators of excitement might be the change in pupil size (i.e. we expect pupil to dilate during excitement and contract otherwise).

There are various tasks and scenes and the level of excitement (pupil dilation) might differ in each:

- in **picture puzzles** we might expect that the level of excitement is kept high during the whole duration, peaking when subject finds a difference
- in ***Where is Waldo?*** puzzle we might expect the level of excitement being low as the subject is having difficulties finding Waldo and peaking when (if) Waldo is found
- viewing a **blank scene** might not excite much

Therefore, it is interesting to ask which tasks and scenes are the most stimulating and whether any significant distinction can be made.

Table 1: Sample of the data for hypothesis 1.

	participant_id	trial	fv_fixation	task_type	stimulus_id	time	x	y	rp
0	022	001	FreeViewing	Puzzle	puz008	0	454.42	317.175	585
1	022	001	FreeViewing	Puzzle	puz008	2	454.66	312.000	585
2	022	001	FreeViewing	Puzzle	puz008	4	453.30	312.375	585
3	022	001	FreeViewing	Puzzle	puz008	6	454.10	309.450	587
4	022	001	FreeViewing	Puzzle	puz008	8	454.82	310.725	585

Hypothesis 1: Saccade Frequency

We will restrict ourselves only to the following hypothesis:

- H_0 (*null hypothesis*) : The frequency of saccades does not decrease when viewing a scene.
- H_1 (*alternative hypothesis*) : The frequency of saccades decreases at some point in the trial.

Load Data

We will restrict ourselves only to one subject. We have arbitrarily chosen the participant with ID 022. We select and load all *Free Viewing* data with **Natural** and **Blank** scenes. We will be interested in columns:

- **time** : Indicating the time in milliseconds from the beginning of the trail.
- **x, y** : The position on the screen where the participant was looking at the given time.
- **rp**: Pupil size of the right eye.

In Table 1 we provide a sample of the data we will use.

Data Manipulation

First step of the pre-processing will be to remove the blinks. We look at the pupil size and consider every value smaller than some threshold to be an indication of a blink. The average pupil size to be between 1 and 8 mm (“Pupil Diameter,” n.d.). We set 1 mm as our threshold. We will then discard a portion of our data around such points. The average blink duration ranges between 100 and 400 ms (“Average Duration of a Single Eye Blink,” n.d.). We delete another 200 ms before and after each blink/semi-blink to eliminate the initial and final parts in which the pupil was still partially occluded.

For each trial, we will detect saccades using the [REMoDNaV](#) python library (Dar, Wagner, and Hanke 2019). This gives us the start and end times of saccades (and other events) as

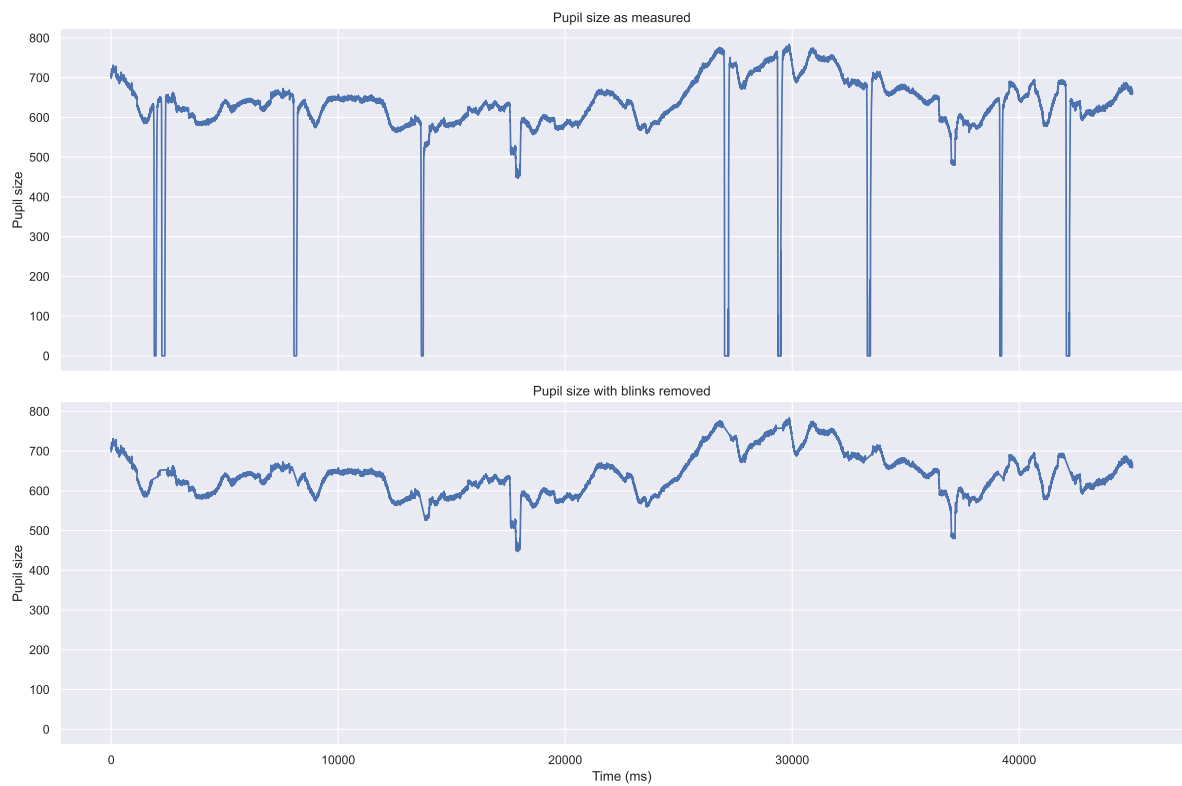


Figure 1: Pupil size data with blinks (top) and with blinks removed (bottom).

Table 2: Sample of Freeviewing data.

	participant_id	trial	fv_fixation	task_type	stimulus_id	start_time
0	022	003	FreeViewing	Natural	nat010	258
1	022	003	FreeViewing	Natural	nat010	480
2	022	003	FreeViewing	Natural	nat010	670
3	022	003	FreeViewing	Natural	nat010	1166
4	022	003	FreeViewing	Natural	nat010	1266

well as additional information such as start and end position of the gaze or peak velocity. We will consider only the start times of the saccades and disregard all other data. Table 2 lists a sample of the processed data.

Methodology and Exploration

To help us understand how we should model the observations, we first look at the distribution of saccades in time. The x axis represents the time from the start of a trial and each row is a sequence of dots representing the start time of a saccade.

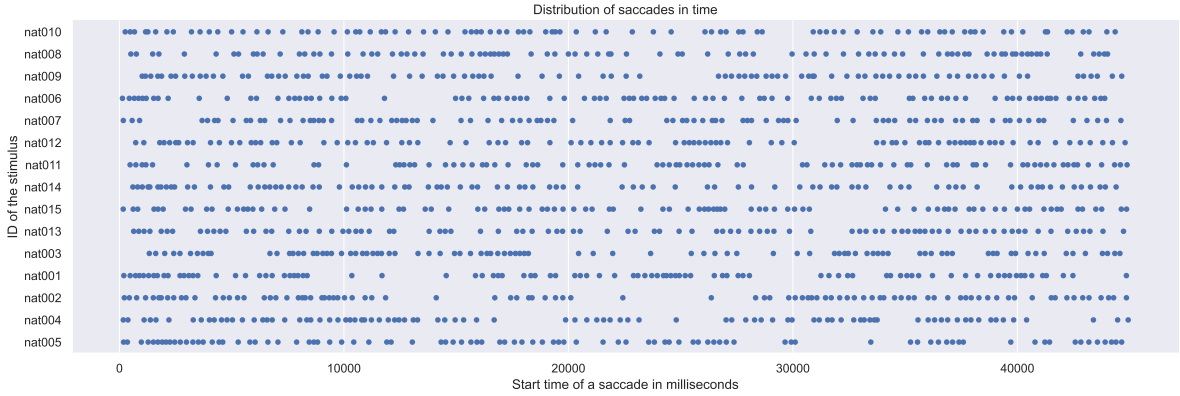


Figure 2: Distribution of saccades in time in Natural scenes for one participant.

As the occurrence of a saccade is a complex process that depends on the scene viewed, mental state of the participant, tiredness or maybe even things like personal history, eyesight, mood and many other factors, it is impossible for us to accurately predict the occurrence of the next saccade. To simplify the situation, we might look at it as on a stochastic process, where for some small enough time interval, we have a fixed probability for an occurrence of a saccade.

Let N_t be a variable counting the number of saccades in a time interval $[0, t]$. Our assumption can be formulated as

$$\lim_{t \rightarrow 0} \Pr[N_t \geq 1]/t = \lambda$$

for some $\lambda \in \mathbb{R}^+$. That is, the probability of an saccade in a small time interval $[0, t]$ tends to λt .

Due to physical human limitations, we also have

$$\lim_{t \rightarrow 0} \Pr[N_t \geq 2]/t = 0$$

or said differently, for a small enough time interval, it is impossible to have two or more successive saccades.

The number of saccades in some interval $[t, s]$ can be expressed then as $N(t) - N(s)$. Additionally, we will assume that for any two disjoint time intervals $[t_1, t_2]$ and $[t_3, t_4]$, the distribution of $N(t_2) - N(t_1)$ is independent of the distribution of $N(t_4) - N(t_3)$.

With these assumptions in place, we are justified to model the occurrences of saccades as a *Poisson process* with rate λ (Mitzenmacher and Upfal 2012).

The theory of Poisson processes then tells us that the number of saccades in some time period of length t is a Poisson random variable with expectation λt . We will divide the time of the experiment to seconds (so $t = 1$) and count the number of saccades in each second. The Figure 3 illustrates such counts for one chosen trial.

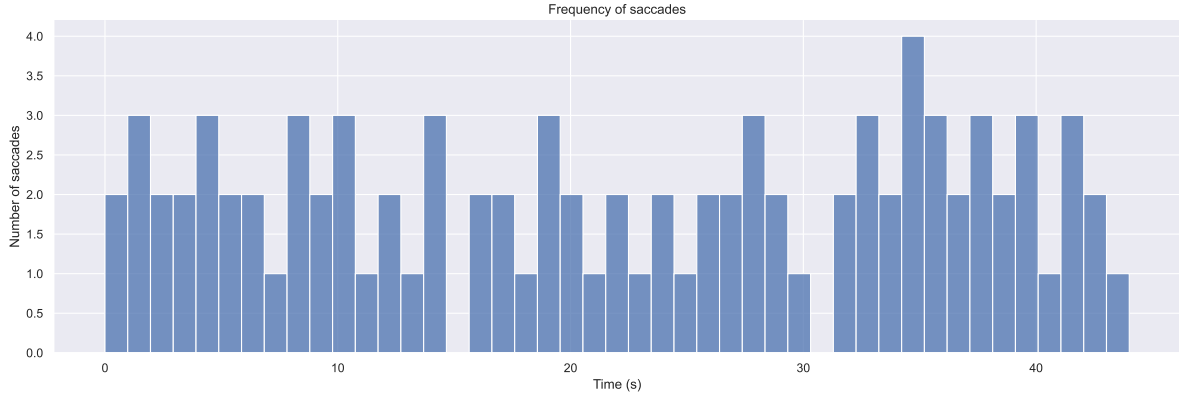


Figure 3: Frequency of saccades in time for one chosen natural scene.

We are looking for a change in behaviour in our participants. To model it we will assume that in some initial time interval the number of saccades will follow some Poisson process with rate λ_1 up until some time τ after which *something* changes and the saccades will follow a different process with rate λ_2 .

Let S_t denote the count of saccades at time t . We have

$$S_t \sim \text{Poisson}(\lambda^{(t)})$$

However, we do not know what the parameter $\lambda^{(t)}$ is. It might be either λ_1 or λ_2 depending on where we put the changing point τ , so

```

\lambda^{\{t\}} =
\begin{cases}
\lambda_1 & t < \tau \\
\lambda_2 & t \geq \tau
\end{cases}

```

We will try to infer these variables from our data using Bayesian statistics. We will need to choose prior distributions to do that. For λ_1 and λ_2 we need to choose from continuous distribution that is able to obtain (potentially) all positive values. Exponential distribution satisfies this requirement, so we set

```

\begin{align*}
\lambda_1 &\sim \mathrm{Exp}(\alpha) \\
\lambda_2 &\sim \mathrm{Exp}(\alpha)
\end{align*}

```

where α is some hyper-parameter. As we have no prior knowledge about τ , we simply choose

$$\tau \sim \mathrm{Unif}(0, 44)$$

To help with the inference, we set the hyperparameter α such that

$$\frac{1}{\alpha} = \mathbb{E}[\lambda \mid \alpha] \approx \frac{1}{45} \sum_{t=0}^{44} S_t$$

which is *about* the value we would expect the λ s to be distributed around.

i Note

This approach was adapted from the first chapter of [Bayesian Methods for Hackers](#).

For our modelling we will use the [PyMC](#) (Salvatier, Wiecki, and Fonnesbeck 2016) implementation of Monte Carlo Markov Chains, using the NUTS sampler with default parameters to generate 5000 samples from the posterior distribution.

Figure 4 shows the posterior distributions of the parameters τ , λ_1 , λ_2 and $\Delta = \lambda_2 - \lambda_1$, while Table 3 provides summary statistics from the chain.

Notice that the posterior distribution of Δ is centered around 0, leading us to believe there is no change in the the rate of saccades. In case there was a change, we would expect a shift either to the positive or negative side. This notion is further reinforced by looking at the Highest Density Interval (HDI) of the posterior distribution of τ which basically spans the entire duration of the trial.

We have




Figure 4: The posterior distribution of the model parameters.

Table 3: Summary statistics for the example chain.

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
tau	23.924	15.010	1.000	44.000	0.549	0.389	669.0	1170.0	1.00
lambda_1	2.035	0.476	1.126	3.035	0.036	0.026	261.0	131.0	1.01
lambda_2	1.993	0.743	0.352	3.027	0.013	0.009	3067.0	1443.0	1.00
delta	-0.042	0.893	-1.791	1.258	0.032	0.022	627.0	953.0	1.01

$$\Pr[H_0] = \Pr[\lambda_2 \geq \lambda_1] = 0.487$$

And as the probability is quite high, we do not reject the null hypothesis in this example.

Natural scenes

We do the same analysis as above for all trials using *Natural* scenes for subject 022.

Figure 5 shows the posterior density estimate for Δ for all *Natural* scene trials.

There does seem to be some variability from scene to scene, however, from the graph we can still see that *on average* the change of the frequencies seems to be centered around zero, even shifted slightly toward the negative side meaning that it is more probable that the frequency of saccades *decreases* slightly as time progresses.

Blank scenes

We now repeat the exact same procedure, but for trials with *Blank* scene.

The situation here is much more diverse in this case. We can see that we can find trials on both ends of the spectrum when considering the change in frequency of saccades. For example, for the trial 094 we have

$$\Pr[\lambda_2 < \lambda_1] = 0.99875$$




Figure 5: Posterior density of Δ for all natural scenes of one chosen subject.




Figure 6: Posterior density of Δ for all blank scenes of one chosen subject.

so there is a very high probability of a *decrease* in the frequency of saccades at some point in the trial. On the other hand, for trial 061 we have

$$\text{Pr}[\lambda_2 > \lambda_1] = 0.9948$$

meaning a high probability of an *increase* in frequency.

This seems quite surprising as one might expect consistent behaviour when viewing the same blank scene multiple times.

Results

Considering all natural scenes as a whole, the probability that the frequency of saccades does not decrease is

$$\text{Pr}[H_0] = \text{Pr}[\lambda_2 \geq \lambda_1] = 0.401$$

which is not low enough to reject the null hypothesis

Considering all the blank trials as a whole, the probability of increase in saccades is

$$\text{Pr}[H_0] = \text{Pr}[\lambda_2 \geq \lambda_1] = 0.42$$

which is not low enough to reject the null hypothesis

Discussion

In this section, we will examine certain assumptions that may have influenced the outcome of our model. One assumption made is that the distribution of saccades is independent of the selected time interval for examination. However, it is possible that the number of saccades may vary over time due to factors such as participant fatigue.

Additionally, we have assumed that there is a single point of change within the data. However, it is feasible that there may be multiple points of change or that the change may occur gradually. Further exploration of these possibilities may serve as a means for model improvement.

Furthermore, our model utilizes specific priors for the parameters. While we have chosen what we believe to be conservative priors, alternative options such as Half-Cauchy or Truncated Normal distributions for the λ s should also be considered. Additionally, the selection of hyperparameters should be evaluated. We claim that the choice of hyperparameters did not

Table 4: Sample of data for hypothesis 2

	participant_id	trial	fv_fixation	task_type	stimulus_id	time	rp
0	022	001	FreeViewing	Puzzle	puz008	0	585
1	022	001	FreeViewing	Puzzle	puz008	2	585
2	022	001	FreeViewing	Puzzle	puz008	4	585
3	022	001	FreeViewing	Puzzle	puz008	6	587
4	022	001	FreeViewing	Puzzle	puz008	8	585

have a significant impact on the model’s results and served primarily to improve the speed of convergence of the chain.

Lastly, we did not take into account the examination of trials as a cohesive unit. Each trial was analyzed individually, but it is possible that the order, type and difficulty of prior trials may have affected performance on subsequent trials.

Hypothesis 2: Pupil dilation

We will study the following hypotheses:

- $H_0^{(T)}$ (*null hypothesis*) : The mean pupil size is equal when comparing Blank scene and a type T scene
- $H_1^{(T)}$ (*alternative hypothesis*) : The mean pupil size is *not* equal when comparing Blank scene and a type T scene

Load Data

Again, we will limit our analysis to data collected from a single participant, specifically participant ID 022. The participant was selected arbitrarily. We have selected and loaded all data collected during the “Free Viewing” task. Our focus will be on the following columns of the dataset:

- **time**: representing the time in milliseconds elapsed from the beginning of the trial.
- **rp**: indicating the size of the right pupil.

In Table 4 we list a sample of data we work with.

Data Manipulation

Again, we remove the blinks and then calculate the mean pupil size for each trial.

Table 5: Average pupil size (right eye) for each task type.

	task_type	avg_rp
0	Puzzle	734.91
1	Waldo	664.90
2	Natural	596.14
3	Blank	554.77
4	Waldo	657.94

Exploration

Looking at Figure 7, we can see that there except for the “Natural” scenes there are no significant outliers in any other category. Table 6 lists summary statistics about the average pupil size.

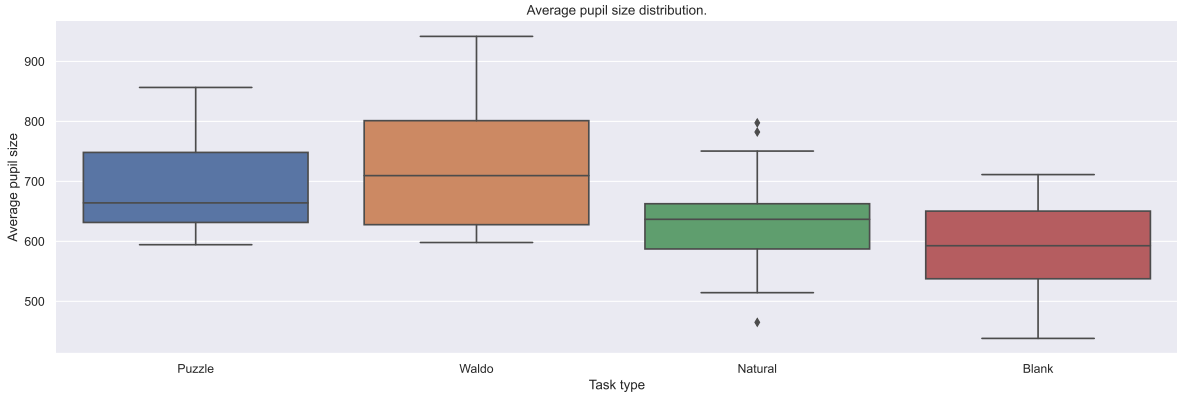


Figure 7: Distribution of the average pupil size across different task types.

Let us look at the distribution of the means. If we want to use t-test, we want our data to be roughly normally distributed.

Looking at the graphs in Figure 8, the distribution of the data seems to be roughly normal, however it is hard to gauge this from such a small sample size. We use the Shapiro-Wilk test, which tests the null hypothesis that the data was drawn from a normal distribution.

Resulting p-values from the Shapiro-Wilk test:

```
Blank    : 0.374
Puzzle   : 0.0556
Waldo    : 0.142
Natural  : 0.592
```

Table 6: Summary statistics for average pupil sizes.

task_type	avg_rp							
	count	mean	std	min	25%	50%	75%	max
Blank	15.0	589.78	89.22	438.20	537.63	592.73	650.43	711.32
Natural	15.0	635.84	91.14	464.97	587.33	636.75	662.73	797.62
Puzzle	15.0	699.44	89.15	594.43	631.53	664.14	748.29	856.61
Waldo	15.0	724.14	112.28	598.10	627.82	709.60	801.23	941.75




Figure 8: Distribution of average pupil sizes by task type.

Using significance level of $\alpha = 0.05$, we conclude that the data is distributed normally (although this is a bit debatable in the case of the “Puzzle” task).

Results

Given that we are evaluating the same participant across different tasks, it is appropriate to use the paired t-test as the data points are linked to one another. This statistical test allows for the comparison of each task with our chosen baseline, which is the blank scene.

Resulting p-values from the paired t-test:

```
Puzzle : 0.00222
Waldo  : 0.00113
Natural : 0.173
```

Thus using significance level of $\alpha = 0.05$, we reject the null hypothesis for the *Puzzle* and *Waldo* tasks, meaning the mean pupil size differs from the baseline (greater as can be seen from Figure 7). We were not able to reject the null hypothesis in case of the *Natural* task.

Discussion

Another interesting metric to consider might be the number of times the pupil size peaks and the count of such peaks.

References

- “Average Duration of a Single Eye Blink.” n.d. *Average Duration of a Single Eye Blink - Human Homo Sapiens - BNID 100706*. <https://bionumbers.hms.harvard.edu/bionumber.aspx?id=100706&ver=4&trm=blink%2B&org=>.
- Dar, Asim H., Adina S. Wagner, and Michael Hanke. 2019. “REMoDNaV: Robust Eye-Movement Classification for Dynamic Stimulation,” April. <https://doi.org/10.1101/619254>.
- Mitzenmacher, Michael, and Eli Upfal. 2012. *Probability and Computing Randomized Algorithms and Probabilistic Analysis*. Cambridge Univ. Press.
- “Pupil Diameter.” n.d. *Pupil Diameter - Human Homo Sapiens - BNID 105349*. <https://bionumbers.hms.harvard.edu/bionumber.aspx?id=105349&ver=4&trm=pupil%2Bsize&org=>.
- Salvatier, John, Thomas V. Wiecki, and Christopher Fonnesbeck. 2016. “Probabilistic Programming in Python Using PyMC3.” *PeerJ Computer Science* 2 (April): e55. <https://doi.org/10.7717/peerj-cs.55>.